

# Foundations of Computer Vision

## Learning Objectives | Outcomes

- What is Computer Vision?
- Understand Basic Concepts in CV.
- Use cases in Healthcare, Security, Entertainment, etc.
- Understand Digital Image Formats (RGB vs Grayscale vs Binary Images).
- Digital Image Structures (Pixel values, Grids, Color Channels, Color Spaces, Image Shapes).
- Load, Display, and save Images using OpenCV.
- Perform Image Manipulations (Resizing, Cropping, Rotation) using OpenCV.
- Apply real-time Filters: Blurring, Edge Detection and etc.
- Build an Interactive Filter app in Python.
- Apply basic Statistics and Visualizations Data Plots & Histograms from Visual Datasets.
- Grasp the impact of preprocessing for ML/DL pipelines.
- Build simple ML models to Classify image data
- Construct and Train beginner-level CNNs using Python frameworks
- Translate theory into working projects using Python, OpenCV, PyTorch, and more.

# Introduction to Computer Vision & Image Types

## Topics Covered:

- What is Computer Vision? Scope and applications.
- Image formats: RGB vs. Grayscale vs. Binary, Image Channels and Pixel Intensity.
- Image Histograms, Formats (JPG, PNG, BMP), Resolution Concepts.

**Tools Used:** Python, OpenCV, Pillow

**Visualization:** Pixel matrix Visualizer, Histogram Plots

## What is Computer Vision?

### Definition:

Computer Vision is a field of Artificial Intelligence that empowers machines to interpret images and videos, enables computers to understand and interpret visual information from the world (Images and Videos), Simulating Human Vision.

## Examples and Definitions of real-world Applications:

- Face Recognition (like unlocking phones)
- Medical Imaging Analysis (detecting tumors)
- Self-Driving | Autonomous Driving Cars (lane and object detection)

## Image Types

- ❖ Color Channels (RGB) vs. Grayscale Images
- ❖ Pixel Matrix Representation (Image as Matrices of Pixels)

### RGB Images:

- Made of 3 channels: Red, Green, Blue
- Each pixel has 3 values (e.g. [255, 0, 0] for red)

### Grayscale Images:

- Only 1 channel, where each pixel holds brightness
- Values range from 0 (black) to 255 (white)

## 2 Image Manipulation with OpenCV

### Tutorial Steps with Code Sample Preview (Python):

- Read, Load and display images in OpenCV (cv2.imread, cv2.imshow)
- Convert Images to Grayscale
- Color Space Conversions (RGB ↔ HSV, Grayscale)
- Inspect Pixel Values and Dimensions
- Resizing, Cropping, Rotating Images (using Transformation Matrix)
- Applying Filters: Blur, Canny Edge Detection (cv2.GaussianBlur, cv2.Canny)
- Drawing Shapes, Text Overlays
- Understand basic Image Transformations and Visual Effects.

#### ➤ Hands-on a Mini Project:

##### Build a Simple Photo Filter GUI App:

Design a Python app that loads an image and lets users apply filters (grayscale, blur, edge detection) via button clicks using **tkinter + OpenCV**.

-  UI with sliders for filter intensity (e.g., brightness, blur, contrast)
-  Use OpenCV to apply effects in real-time and (filter selection)
-  Real-time display using OpenCV **imshow** and callbacks
-  Save filtered images to disk

### Tools & Setup

- **Language:** Python 
- **Dependencies (Libraries):** 

From Bash (Terminal, CMD or Shell) run:    pip install **opencv-python matplotlib numpy**

- **Environment:** Jupyter Notebook or Google Colab (recommended for ease of use)

 Load an Image, Printing Shape & Pixel Dept in a Point, Converting RGB to Grayscale, Displaying the Images.

```
import cv2 # Import OpenCV Library

rgb_img = cv2.imread('photo.jpg') # Loads in BGR Image by default
print(rgb_img.shape) # Show image dimensions
print(rgb_img[100, 150]) # Print Pixel intensity at (100, 150)
gray_img = cv2.cvtColor(rgb_img, cv2.COLOR_BGR2GRAY) # Covert BGR to Grayscale
print(gray_img[100, 150]) # Print Pixel intensity at (100, 150)

cv2.imshow('RGB', rgb_img) # Show RGB Image
cv2.imshow('Grayscale', gray_img) # Show Grayscale Image
cv2.waitKey(0) # Keeping Window Open Until Pressing a Key on the Keyboard
cv2.destroyAllWindows() # Closing All Windows
```

## Visualize an Image in a Plot

```
import cv2
import matplotlib.pyplot as plt # Import Plot Library

image = cv2.imread("photo.jpg")
image_rgb = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
plt.imshow(image_rgb) # Show Plot of the Image
plt.title("Good Photo") # Setting Title for the Plot
plt.axis("off") # Hides both the x-axis and the y-axis of the Plot
plt.show() # Show the Plot
```

 Resize, Crop, Filter, Rotate, Edge Detection, Color Space Conversions, Convert Formats (BGR ↔ HSV)

```
import cv2

img = cv2.imread('sample.jpg')
resized = cv2.resize(img, (100, 100)) # Resized
cropped = img[50:150, 100:200] # Cropped dimension range: y-range, x-range
blurred_filtered = cv2.GaussianBlur(img, (5, 5), 0) # Applied GaussianBlur Filter
rotated = cv2.rotate(img, cv2.ROTATE_90_CLOCKWISE) # Rotate 90 Degree Clockwise
edges = cv2.Canny(img, 100, 200) # Edge Detection by Canny
RGB_to_HSV = cv2.cvtColor(img , cv2.COLOR_RGB2HSV) # Convert rgb to hsv
HSV_to_RGB = cv2.cvtColor(img , cv2.COLOR_HSV2RGB) # Convert hsv to rgb
cv2.imshow("NameOfImage", oneOfRawImageOrProcessedImagesAbove)
cv2.waitKey(0)
```

## 3 Basic Statistics & Data Visualization for Vision Datasets

### Objective:

- Explore pixel data distributions and summarize image information.
- Perform basic statistics on image datasets.
- Visualize class distributions.
- Detect anomalies or imbalance.

### Libraries:

- Numpy
- Matplotlib
- seaborn

### Tutorial Steps:

- ✓ Extract image channels and compute statistics
- ✓ Visualize grayscale histograms
- ✓ Compare pixel intensity distributions across images
- ✓ Understanding Data & Visualize Dataset Stats
- ✓ Dataset shape, summary statistics (mean, variance, std), class distribution
- ✓ Data distributions and outlier detection
- ✓ Explore vision datasets (MNIST, CIFAR)
- ✓ Exploratory Data Analysis (EDA) of image datasets
- ✓ Plotting (visualization) with Matplotlib & Seaborn
- ✓ Histograms and box plots
- ✓ Plot distributions, calculate mean/std, Mean and variance of pixel intensity

**Mini Project:** Visualize class-wise image distribution in MNIST or CIFAR-10

### Image Dataset Analysis

Understanding dataset before training helps identify class imbalance, image anomalies, or corrupted samples.

### Visualize Dataset Distribution

```
import matplotlib.pyplot as plt
import seaborn as sns

labels = ['Cat']*500 + ['Dog']*1000
sns.countplot(x=labels)
plt.title("Class Distribution")
plt.show()
```

## Summary Stats

```
import numpy as np
import cv2

image = cv2.imread("photo.jpg")
gray_img = cv2.cvtColor(image, cv2.COLOR_RGB2GRAY)
mean_pixel = np.mean(gray_img)
std_pixel = np.std(gray_img)
print(f"Mean: {mean_pixel:.2f}, Std: {std_pixel:.2f}")
```

## Histogram Plot

```
import matplotlib.pyplot as plt
import cv2

image = cv2.imread("photo.jpg")
gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
gray_flat = gray.flatten()
plt.hist(gray_flat, bins=50, color='gray')
plt.title("Pixel Intensity Distribution")
plt.xlabel("Intensity")
plt.ylabel("Frequency")
plt.show()
```

## Visualize an Image Matrices in a Plot

```
import cv2
import numpy as np
import matplotlib.pyplot as plt

# Load image
img = cv2.imread("photo.jpg")
gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

# Show image matrices
plt.figure(figsize=(8,4))
plt.subplot(1,2,1); plt.title("RGB"); plt.imshow(cv2.cvtColor(img, cv2.COLOR_BGR2RGB))
plt.subplot(1,2,2); plt.title("Grayscale"); plt.imshow(gray, cmap='gray')
plt.show()
```

## Summary Stats

```
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np

# Simulate labels
labels = ['Cat']*300 + ['Dog']*700
sns.countplot(x=labels)
plt.title("Class Distribution")
plt.show()

# Image stats
from PIL import Image
gray = Image.open("photo.png").convert("L")
pixels = np.array(gray).flatten()
print("Mean:", np.mean(pixels))
print("Std Dev:", np.std(pixels))
```

## 4 Fundamentals of Supervised Machine Learning (ML) for CV

### 💡 Objective:

- ❖ Introduce ML concepts and apply them to visual data features.
- ❖ Learn linear regression and decision trees
- ❖ Train models on image-derived features

### Definition:

Learning from labeled data where each input (image) has a known output (label). The model finds patterns to predict labels on new unseen images.

### • Topics Covered:

- Linear Regression for pixel intensity prediction
  - Decision Trees on image features
  - Extract features from images (pixel intensity, color histograms, mean pixel value, edge count)
  - Create a dataset for classification (e.g., cat vs dog)
  - Train, test, split, cross-validation, evaluation metrics and model (accuracy, precision)
  - Train a decision tree classifier
- Libraries: Scikit-learn, NumPy
  - Code Example: with scikit-learn (sklearn)

### 💡 Key Algorithms:

- LinearRegression
- DecisionTreeClassifier

### 📐 Example Algorithms

#### Linear Regression (for image-derived features):

```
from sklearn.linear_model import LinearRegression

X = [[20], [40], [60]] # Feature: brightness
y = [1, 2, 3]           # Target label

model = LinearRegression()
model.fit(X, y) # model.fit(X_train, y_train)
print("Prediction for brightness=50:", model.predict([[50]]))
print(f'Accuracy: {model.score(X, y):.2f}') # model.score(X_test, y_test)
```

## Decision Trees:

```
from sklearn.tree import DecisionTreeClassifier

X = [[128, 0], [200, 1]] # Features: avg_brightness, edge_flag
y = [0, 1]                 # Classes: Cat=0, Dog=1

clf = DecisionTreeClassifier()
clf.fit(X, y)
print("Class Prediction:", clf.predict([[150, 0]]))
```

## Linear Regression + Decision Trees:

```
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeClassifier

# Toy data
X = [[0], [100], [150]] # brightness
y = [0, 1, 1]             # label

# Linear model
lr = LinearRegression()
lr.fit(X, y)
print("Prediction for 120:", lr.predict([[120]]))

# Decision Tree
dt = DecisionTreeClassifier()
dt.fit(X, y)
print("Tree Prediction:", dt.predict([[120]]))
```

## 5 Intro to Convolutional Neural Network (CNN): Architecture & Use Cases

### Objective:

- ❖ Understand basic structure of CNNs layers and their applications.
- ❖ Understand MNIST digit classifier with PyTorch
- ❖ Train a classifier on MNIST

### What is a CNN?

#### Definition:

A Convolutional Neural Network (CNN) is a specialized deep learning model for image data. It uses convolution filters to detect patterns like edges, shapes, and textures.

- **Topics Covered:**
  - What are Convolutions, Pooling, FC layers?
  - Activation functions, Flattening layers, Kernel, Stride, Padding
  - CNN layers: Conv2D, MaxPooling, Fully Connected (FC)
  - Popular architectures: LeNet, AlexNet (intro-level)
  - Feed-forward architecture for image classification
  - Evaluation metrics: accuracy, confusion matrix
  - Building a basic CNNs with Keras or PyTorch
  - Train on MNIST dataset
  - Visualize predictions and misclassifications
- **Libraries:** MNIST digit classifier using PyTorch (Torchvision, Torchvision.Datasets.MNIST)
- **Code Snippet:**

```
class SimpleCNN(nn.Module):  
    def __init__(self):  
        super().__init__()  
        self.conv1 = nn.Conv2d(1, 16, 3)  
        self.fc = nn.Linear(16*26*26, 10)  
  
    def forward(self, x):  
        x = F.relu(self.conv1(x))  
        x = x.view(x.size(0), -1)  
        return self.fc(x)
```

### Common Layers Explained:

- **Conv2D:** Applies filters to extract features
- **ReLU:** Introduces non-linearity
- **MaxPooling (MaxPool2D):** Reduces image size while preserving info (DownSampling)
- **Fully Connected:** Final decision-making
- **Linear:** Classification

## 34 MNIST Digit Classifier (Classify Handwritten Digits by MNIST)

- Load dataset using `torchvision` or `keras.datasets` (Load MNIST data and preprocess)
- Build and Train a simple CNN with 2–3 layers
- Evaluate accuracy and visualize predictions
- Plot misclassified digits and confusion matrix

### 🛠 Hands-On Projects

#### 💡 Example: MNIST Digit Classifier

```
import torch.nn as nn
import torch.nn.functional as F

class SimpleCNN(nn.Module):
    def __init__(self):
        super(SimpleCNN, self).__init__()
        self.conv1 = nn.Conv2d(1, 16, 3) # 1-channel grayscale input
        self.pool = nn.MaxPool2d(2, 2)
        self.fc1 = nn.Linear(16*13*13, 10) # Output 10 digits

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = x.view(-1, 16*13*13)
        x = self.fc1(x)
        return x
```

Use with PyTorch's `torchvision.datasets.MNIST` for training.

#### 💻 Sample Code (Keras)

```
from keras import models, layers, datasets
(x_train, y_train), (x_test, y_test) = datasets.mnist.load_data()
x_train = x_train[..., None] / 255.0
x_test = x_test[..., None] / 255.0
model = models.Sequential([
    layers.Conv2D(32, (3, 3),
                 activation='relu',
                 input_shape=(28,28,1)),
    layers.MaxPooling2D((2, 2)),
    layers.Flatten(),
    layers.Dense(100, activation='relu'),
    layers.Dense(10, activation='softmax')
])
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
model.fit(x_train, y_train, epochs=5, validation_split=0.1)
```

## Hands-on Projects

Project	Tools	Description
Photo Filter App	OpenCV + tkinter	GUI for image filtering
MNIST Classifier	PyTorch	Train CNN to classify digits
Class Visualizer	Matplotlib	EDA on dataset distribution
Image Feature Predictor	Scikit-learn	ML with handcrafted features

## Tools & Dependencies

- Python ≥ 3.8
- OpenCV (`opencv-python`)
- PyTorch + torchvision
- Matplotlib + Seaborn
- Scikit-learn

## ML Algorithms & CNNs

- Train a Decision Tree on feature-extracted images
- Build & train a CNN for MNIST
- 🌟 Challenge: “Replace ReLU with LeakyReLU and compare results”

```
import torch.nn as nn
class Net(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv = nn.Conv2d(1, 16, 3)
        self.fc = nn.Linear(16*26*26, 10)

    def forward(self, x):
        x = F.relu(self.conv(x))
        x = x.view(x.size(0), -1)
        return self.fc(x)
```

## Why Preprocess?

- Ensuring consistent input for CNNs
- Highlighting meaningful features

## Mini Project: Build Your Own Image Editor

Create a notebook where students upload a photo, apply transformations, and save the result. Bonus challenge: add a GUI with `cv2.imshow()` + keyboard controls!

## Interactive Tasks

1. Crop the image to center the face.
2. Apply different kernel sizes for blurring.
3. Build a tkinter GUI to toggle filters.

## Suggested Projects (Progressively Challenging)

Project	Skills Practiced	Tools
MNIST Classifier	CNN basics	PyTorch, torchvision
Dog vs. Cat Classifier	Data pipelines, fine-tuning	PyTorch, transfer learning
Custom YOLOv5 Detector	Object detection	Ultralytics YOLO repo
Semantic Segmentation	Masking, pixel-wise classification	PyTorch + OpenCV

## Crystal-Clear Documentation Style

Module Doc Example

## Visual Flow Diagram

[Insert block diagram showing image flow from file → preprocessing → display]

## Extra Code Examples and Notebooks

- [x] CNN training with PyTorch (starter + intermediate)
- [x] Object detection pipeline with pre-trained YOLOv5
- [x] Semantic segmentation using U-Net
- [x] Jupyter notebook: Data Augmentation walkthrough
- [x] Custom dataset loader tutorial (images + annotations)